

Project tutoring

Alessandro Di Federico

ale@clearmind.me

Algorithms and Parallel Computing
Mathematical Engineering
Politecnico di Milano

December 17, 2015

Index

Dividing your program

Data structures 101

Linked lists

How to divide your program

- Identify the core components:
 - a set of classes/functions closely related
 - keep them together if they interact frequently
 - split them if they have a tiny interface
- Create a .cpp file and .h file for each component

What to put in .h files

- In general:
 - `.cpp` files contain definitions
 - `.h` files contain declarations
- Suppose a `.cpp` file needs a function defined elsewhere
- It must include a `.h` file containing its declaration

Example

```
// example.h
int factorial(int i);

// example.cpp
#include "example.h"

int factorial(int i) {
    if (i > 1)
        return i * factorial(i - 1);
    else
        return i;
}
```

Example

```
// user.cpp
#include "example.h"

int main() {
    factorial(4);
    return 0;
}
```

Static functions

- Suppose a function is used only in a single .cpp
- Define it there, no declaration in .h files
- Mark it `static`

Static function example

```
// user.cpp
#include "example.h"

int main() {
    factorial(4);
    return 0;
}
```


Inline functions

- In general .h files should only contain declarations
- Suppose you put a definition in an .h file
- Suppose the .h file is included by multiple .cpp files
- The compiler will find multiple definitions and will complain
- If you really want to put a definition in an .h file, use `inline`

Bad

```
// example.h  
int meaning() { return 42; }
```

```
// example.cpp  
#include "example.h"
```

```
int main() {  
    meaning();  
}
```

OK

```
// example.h  
inline int meaning() { return 42; }
```

```
// example.cpp  
#include "example.h"
```

```
int main() {  
    meaning();  
}
```

Classes

- If a class is used by a single .cpp file it can stay there
- If it's used by multiple .cpp files put it in a .h file

Method definitions

- Two options:
 - define it inside the class (*inline* definition)
 - define it outside the class (*outline* definition)
- It's OK to put inline definitions in .h files
- They are implicitly inline functions

Method definition examples

```
// example.h
class MyClass {
    int methodInline() { return 42; }
    int methodOutline();
};
```

```
// example.cpp
#include "example.h"

int MyClass::methodOutline() {
    return 44;
}
```

And with CMake?

Just add to your executable all your .cpp files

```
add_executable(hello hello.cpp  
               example.cpp  
               car.cpp)
```

Index

Dividing your program

Data structures 101

Linked lists

Dense matrix

- What's the best way to represent a dense matrix?
- A sequential memory area of $M \times N$ elements

Size known statically

```
const int row_count = 1234;  
const int col_count = 4321;  
  
int denseMatrix[row_count][col_count];
```

Size not known statically

```
int *denseMatrix = new int[row_count * col_count];  
// A std::vector is fine too  
  
// To access element (x, y)  
denseMatrix[x * col_count + y];
```

Dense matrix: pros & cons

- Suppose we have T non-zero elements:
 - Memory taken: $M \times N$ elements
 - Access time: constant
 - Enumeration time: $M \times N$

Sparse matrix

- Assume the full matrix won't fit in memory
- We can represent only the non-zero elements
- There are various options
- e.g. an unordered vector of $\langle x, y, value \rangle$ triples

Dense matrix: pros & cons

- Suppose we have T non-zero elements:
 - Memory taken: T elements
 - Access time: T
 - Enumeration time: T

In general

- Think about your data structure in terms of:
 - memory taken
 - access time
 - enumeration time
- Evaluate it in terms of corner cases:
 - How does it perform with a very sparse matrix?
 - How does it perform with an average matrix?
 - How does it perform with a very dense matrix?
- Try to get good metrics in all cases

Don't use both

- Using a sparse and dense representation is a bad idea
- You get all the pros and cons
In particular you're using a sparse representation but if the input is too large your program won't work anyway
- So don't use both (at the same time)
- You can choose at run-time which representation to use!

Allocations

- Allocations and memory releases take time
- How frequently do you allocate/release memory?
- Our problem is pretty static, after initialization:
 - grid size is constant
 - total amount of cars is constant

Index

Dividing your program

Data structures 101

Other suggestions

Optimizations

- To test the performance of your code, enable optimizations!
- Pass the “-O2” flag to the compiler
- Or simply build in release mode with CMake¹

¹See CMake slides

Compiler warnings

- Compiler warnings might seem boring
- 99% of the times it's a sign of something wrong
- Take the time to fix them all
- You can get additional warnings with the “-Wall” flag
- Don't send me code which triggers warning

No magic numbers

- Don't put wild numbers in your code
- The same number might have different meanings
- Always define a global constant

using

- using it's similar to typedef
- Useful to define shorter alias of an existing type

```
using IntVect = std::vector<int>;  
IntVect MyVector;
```

using namespace

```
using namespace something;
```

- Makes something:: implicit

```
using namespace std;  
vector<int> MyVector;  
cout << "Hi!" << endl;
```

- **Never** put using namespace in a .h file

Inheritance or templates?

- Should BlueCar inherit a base class Car?
- Should it specialize a template class Car?

Inheritance

```
class Car {  
    virtual void advance() = 0;  
    // ...  
};  
  
class BlueCar : Car {  
    void advance() {  
        // BlueCar-specific logic  
    }  
    // ...  
};
```

Template specialization

```
const int BlueColor = 1;
```

```
template<int Color>  
class Car {  
    void advance();  
    // ...  
};
```

```
template<>  
class Car<BlueColor> {  
    void advance() {  
        // BlueCar-specific logic  
    }  
    // ...  
};
```

Use inheritance if...

you want to use BlueCar and RedCar interchangeably
through their base class (Car)

```
int myFunction(Car *MyCar) {  
    MyCar->advance();  
}
```

Note: virtual methods incur in a small extra cost upon call.

Use templates if...

you always know if the object you're dealing with it's a
BlueCar or RedCar

```
int myFunction() {  
    std::vector<Car<BlueColor>> BlueCars;  
    // ...  
    for (Car<BlueColor>& MyCar : BlueCars) {  
        // ...  
    }  
  
    std::vector<Car<RedColor>> RedCars;  
    // ...  
}
```

If your laptop sucks

- Some laptops don't support VT-x
- So, the VM will see only a single core
- No real speedups from OpenMP/MPI!

Possible solutions

- Buy a decent laptop
- Install Linux side-by-side on your laptop
- Get in touch with the “Servizio prestito PC”²
- Ask us for access to a server

²Istituto per il diritto allo Studio Universitario (ISU) - Servizio Prestito Personal Computer, Centro Servizi Via Golgi, 20

License



This work is licensed under the Creative Commons Attribution-ShareAlike 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/3.0/> or send a letter to Creative Commons, 444 Castro Street, Suite 900, Mountain View, California, 94041, USA.